

Short Questions to Analyzing the NYC Subway Dataset

Analyzing the NYC Subway Dataset

Short Questions

Overview

This project consists of two parts. In Part 1 of the project, you should have completed the questions in Problem Sets 2, 3, 4, and 5 in the Introduction to Data Science course.

This document addresses part 2 of the project. Please use this document as a template and answer the following questions to explain your reasoning and conclusion behind your work in the problem sets. You will attach a document with your answers to these questions as part of your final project submission.

Section 1. Statistical Test

1.1 Which statistical test did you use to analyze the NYC subway data? Did you use a one-tail or a two-tail P value? What is the null hypothesis? What is your p-critical value?

Answer:

- *I used Mann-Whitney U test to analyze the ridership from NYC subway dataset.*
- *The Python – `scipy.stats.mannwhitneyu` function gives a one sided p value as the result. Hence, I used one sided p-value.*
- *The null hypothesis is – The subway ridership is same on rainy days and on non-rainy days – i.e., the number of people riding the subway when it rains is same as when it doesn't rain.*
- *I used a p-critical value of 0.05.*

1.2 Why is this statistical test applicable to the dataset? In particular, consider the assumptions that the test is making about the distribution of ridership in the two samples.

Answer:

- *Mann-Whitney U test is applicable for non-parametric data i.e., data which is not drawn from any parametric distribution. In case of NYC dataset, the dataset does not come from normal distribution. Hence, Mann-Whitney U test is applicable.*
- *Also, the following assumptions for the NYC subway data hold true which are necessary for conducting Mann-Whitney U Test:*
 - *Observations from both groups are independent of each other*
 - *The dependent variable is either of continuous or ordinal type*
 - *One of the independent variables is of categorical type with two groups*
 - *In order to compare the means of two different groups, the distribution of scores for these two groups should have the same shape*

1.3 What results did you get from this statistical test? These should include the following numerical values: p-values, as well as the means for each of the two samples under test.

Answer:

- *Mean on rainy days: 1105.4463767458733*

- Mean on non-rainy days: 1090.278780151855
- Mann Whitney U Test statistic : 1924409167.0
- P-value: 024999912793489721

1.4 What is the significance and interpretation of these results?

Answer: The Mann Whitney U statistical test resulted in a one sided p-value of 0.025. This p-value is less than the critical value of 0.05 and we can reject the null hypothesis. Hence, we can conclude that the ridership on rainy days is significantly different from that on the non-rainy days.

Section 2. Linear Regression

2.1 What approach did you use to compute the coefficients theta and produce prediction for ENTRIESn_hourly in your regression model:

1. Gradient descent (as implemented in exercise 3.5)
2. OLS using Statsmodels
3. Or something different?

Answer: I used Stochastic gradient descent method. In this method we used randomly selected training samples for calculating new set of parameter values.

2.2 What features (input variables) did you use in your model? Did you use any dummy variables as part of your features?

Answer: I used the following features: rain, precipi, Hour, meantempi, meanpressurei, fog, week of the day.

Yes, I used the dummy features. I converted week of day to dummy_days before using them.

2.3 Why did you select these features in your model? We are looking for specific reasons that lead you to believe that

the selected features will contribute to the predictive power of your model.

- Your reasons might be based on intuition. For example, response for fog might be: “I decided to use fog because I thought that when it is very foggy outside people might decide to use the subway more often.”
- Your reasons might also be based on data exploration and experimentation, for example: “I used feature X because as soon as I included it in my model, it drastically improved my R2 value.”

Answer: I used rain, meantempi, meanpressurei and fog because these are indicators of the weather and weather seems to have an impact on the number of people taking the subway.

I also used 'Hour' and 'Day of the week' as features because these two features capture the routine of the people taking the subway. For eg, during peak hours the subway ridership may be high and the ridership will be different on weekends compared to weekdays.

2.4 What are the coefficients (or weights) of the non-dummy features in your linear regression model?

Answer:

Following are the coefficients of the non dummy features.

Intercept Term: 38.90404671

Rain: -22.47225364,

Precip: 466.2643544

Hour: -50.10849709

Meantempi: -37.24624601,

Meanpressurei: 41.46483278,

Fog: -9.20433213

Day of the week 0: 66.06534061,

Day of the week 1: 93.59672843,

Day of the week 2: 64.4699832 ,

Day of the week 3: 64.93100007,

Day of the week 4: -86.94618904,

Day of the week 5: -173.51590796,

Day of the week 6: 110.85532737

2.5 What is your model's R2 (coefficients of determination) value?

Answer: R2 values: 0.4634

2.6 What does this R2 value mean for the goodness of fit for your regression model? Do you think this linear model to predict ridership is appropriate for this dataset, given this R2 value?

Answer: An R2 value of 0.4613 means using the current set of features used to calculate the value of theta, 46% of the variation in the output can be explained. This also means that there is lot more variation that is not being captured by the features that are used.

This model cannot be used if precise prediction is required. However, in order to understand and predict the trend, this model should be sufficient.


Section 3. Visualization

Please include two visualizations that show the relationships between two or more variables in the NYC subway data.

Remember to add appropriate titles and axes labels to your plots. Also, please add a short description below each figure commenting on the key insights depicted in the figure.

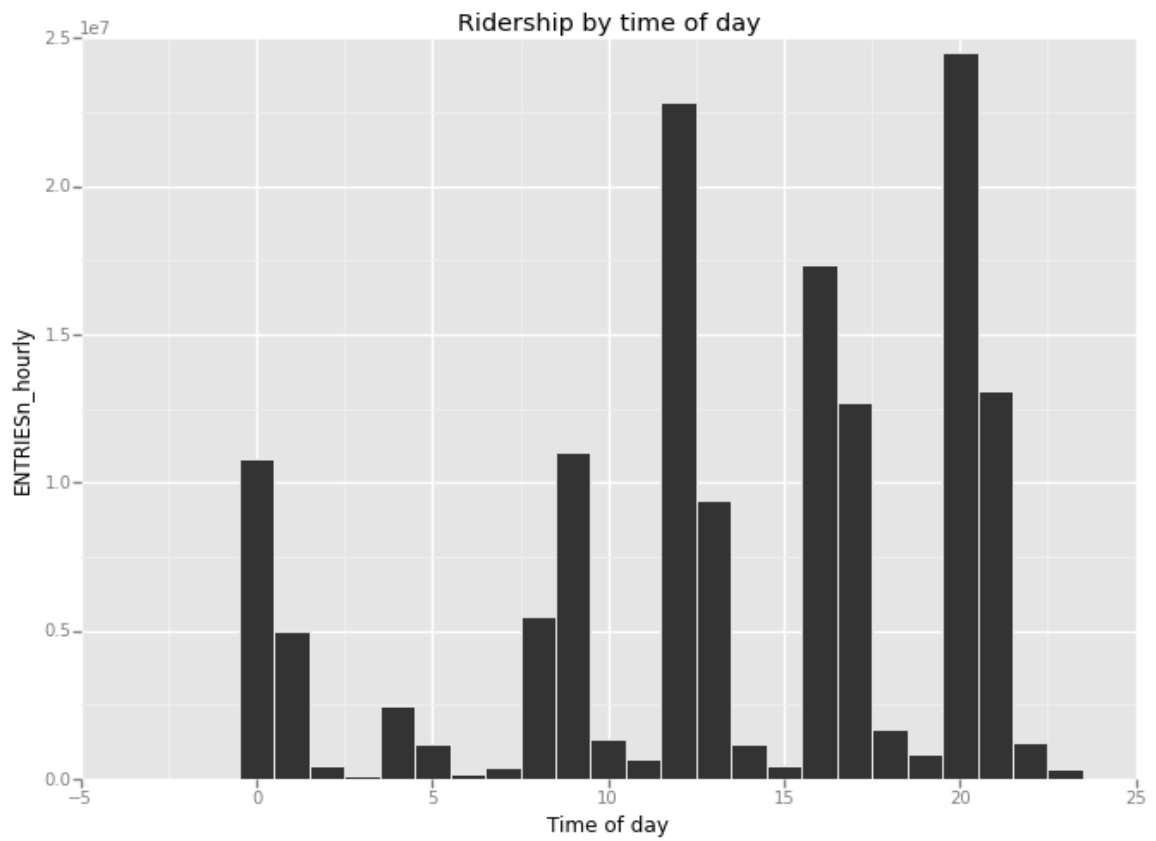
3.1 One visualization should contain two histograms: one of ENTRIESn_hourly for rainy days and one of ENTRIESn_hourly for non-rainy days.

- You can combine the two histograms in a single plot or you can use two separate plots.
- If you decide to use two separate plots for the two histograms, please ensure that the x-axis limits for both of the plots are identical. It is much easier to compare the two in that case.
- For the histograms, you should have intervals representing the volume of ridership (value of `ENTRIESn_hourly`) on the x-axis and the frequency of occurrence on the y-axis. For example, each interval (along the x-axis), the height of the bar for this interval will represent the number of records (rows in our data) that have `ENTRIESn_hourly` that falls in this interval.
- Remember to increase the number of bins in the histogram (by having larger number of bars). The default bin width is not sufficient to capture the variability in the two samples.




3.2 One visualization can be more freeform. You should feel free to implement something that we discussed in class (e.g., scatter plots, line plots) or attempt to implement something more advanced if you'd like. Some suggestions are:


- Ridership by time-of-day




- Ridership by day-of-week



- Stations with most entries during time of day



- Ridership at different stations



Section 4. Conclusion

Please address the following questions in detail. Your answers should be 1-2 paragraphs long.

4.1 From your analysis and interpretation of the data, do more people ride the NYC subway when it is raining or when it is not raining?

Answer: After looking at the statistical tests performed on the subway dataset, we can conclude that more people ride the subway when it is raining.

4.2 What analyses lead you to this conclusion? You should use results from both your statistical tests and your linear regression to support your analysis.

Answer: After conducting the Mann Whitney U test on the NYC subway dataset, we can conclude that the ridership during rainy days is different from that on the non-rainy days. Since the mean ridership on rainy days is greater than that on the non-rainy days we can conclude that more people ride the subway when it is raining.

Also, when we perform linear regression to predict the hourly entries, using rain as a feature resulted in an increased R square value of 0.4217 compared to 0.4156 when rain is not used. This explains that using rain as a feature increases the amount of variation predicted by the independent variables.

Section 5. Reflection

Please address the following questions in detail. Your answers should be 1-2 paragraphs long.

5.1 Please discuss potential shortcomings of the methods of your analysis, including:

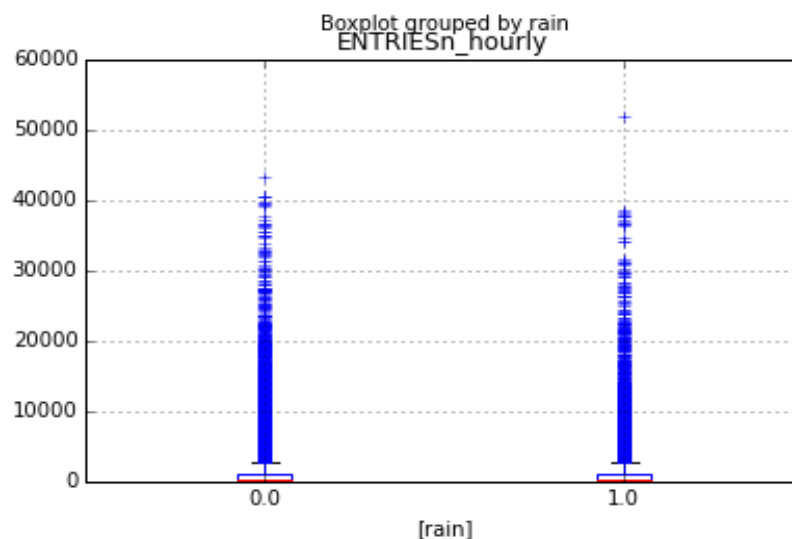
1. Dataset,
2. Analysis, such as the linear regression model or statistical test.

Answer:

- *Considering some of the independent variables like rain, precipi, fog, day of the week, time during the day, etc, most of these variables are binary variables. Logistic regression would be a better method to apply to such variables because the cost function used in linear regression does not suit for binary predictors.*
- *While we perform linear regression on the given dataset, we have to keep in mind the assumptions that need to hold good before we apply linear regression:*
 1. *Linear relationship between the dependent and independent variables*
 2. *Normality of data*
 3. *Homogeneity of data*
 4. *Independence*
- *Some of these assumptions don't hold good incase of our dataset. For instance, we are using variables like rain and mean temperature which are not independent. That is, there is some correlation between these two variables. This may violate the assumption of independence.*

It is difficult to establish the normality of the data in this case. And even if the data is normally distributed, the variance for different variables is not the same. Hence, homogeneity of data does not hold good in many cases.

- After looking at the scatter plots between the dependent variable and different independent variables, it does not seem that some of these variables share a linear relationship with the dependent variable. Hence, linear regression may not be suitable for the current dataset.
- Another important shortcoming of linear regression is that it is very sensitive to outliers. It can be seen from the box plot below that there are many outliers for `ENTRIESn_hourly` in case of rainy and non-rainy days. These outliers can cause incorrect fitting of the data.



5.2 (Optional) Do you have any other insight about the dataset that you would like to share with us?

Published by [Google Drive-Report Abuse](#)–Updated automatically every 5 minutes

Links or references used:

<https://ggplot.yhathq.com/>

<http://web.stanford.edu/~cengel/cgi-bin/anthrospace/ggplot-from-python-with-rpy2>

http://folk.uio.no/rnymoen/ECON4160_v09_Lect5v9.pdf

https://stat.duke.edu/courses/Fall00/sta103/lecture_notes/multregr.pdf

<http://www.ftpress.com/articles/article.aspx?p=2248639&seqNum=5>

http://www.springer.com/cda/content/document/cda_downloaddocument/9780387874579-c2.pdf?SGWID=0-0-45-721017-p173851911.